

Q1. Write a Python program to perform a Z-test for comparing a sample mean to a known population mean and interpret the results.

Ans:-

```
import numpy as np
import scipy.stats as stats

# Function to perform a one-sample Z-test
def z_test(sample, population_mean, population_std, alpha=0.05):
    # Calculate sample mean and sample size
    sample_mean = np.mean(sample)
    n = len(sample)

    # Calculate the Z-score
    z_score = (sample_mean - population_mean) / (population_std /
np.sqrt(n))

    # Calculate the p-value (two-tailed)
    p_value = 2 * (1 - stats.norm.cdf(abs(z_score)))

    # Print results
    print(f"Sample Mean: {sample_mean:.2f}")
    print(f"Population Mean: {population_mean:.2f}")
    print(f"Z-Score: {z_score:.2f}")
    print(f"P-Value: {p_value:.4f}")

    # Interpret the results
    if p_value < alpha:
        print("Reject the null hypothesis: There is a significant
difference between the sample mean and the population mean.")
    else:
        print("Fail to reject the null hypothesis: There is no
significant difference between the sample mean and the population
mean.")

# Example usage
if __name__ == "__main__":
    # Sample data
    sample_data = [20, 22, 19, 24, 21, 23, 20, 22, 19, 21]

    # Known population parameters
    population_mean = 21 # Known population mean
    population_std = 2   # Known population standard deviation

    # Perform the Z-test
    z_test(sample_data, population_mean, population_std)
```

```
Sample Mean: 21.10
Population Mean: 21.00
```

Z-Score: 0.16
P-Value: 0.8744
Fail to reject the null hypothesis: There is no significant difference between the sample mean and the population mean.

Q2.Simulate random data to perform hypothesis testing and calculate the corresponding P-value using Python?

Ans:-

```
import numpy as np
import scipy.stats as stats
import matplotlib.pyplot as plt

# Function to perform a one-sample Z-test
def z_test(sample, population_mean, population_std, alpha=0.05):
    # Calculate sample mean and sample size
    sample_mean = np.mean(sample)
    n = len(sample)

    # Calculate the Z-score
    z_score = (sample_mean - population_mean) / (population_std / np.sqrt(n))

    # Calculate the p-value (two-tailed)
    p_value = 2 * (1 - stats.norm.cdf(abs(z_score)))

    # Print results
    print(f"Sample Mean: {sample_mean:.2f}")
    print(f"Population Mean: {population_mean:.2f}")
    print(f"Z-Score: {z_score:.2f}")
    print(f"P-Value: {p_value:.4f}")

    # Interpret the results
    if p_value < alpha:
        print("Reject the null hypothesis: There is a significant difference between the sample mean and the population mean.")
    else:
        print("Fail to reject the null hypothesis: There is no significant difference between the sample mean and the population mean.")

# Simulate random data
def simulate_data(sample_size, population_mean, population_std):
    # Generate random sample data from a normal distribution
    sample_data = np.random.normal(loc=population_mean, scale=population_std, size=sample_size)
    return sample_data

# Example usage
```

```

if __name__ == "__main__":
    # Parameters for simulation
    sample_size = 30 # Size of the sample
    population_mean = 100 # Known population mean
    population_std = 15 # Known population standard deviation

    # Simulate random data
    sample_data = simulate_data(sample_size, population_mean,
                                population_std)

    # Perform the Z-test
    z_test(sample_data, population_mean, population_std)

    # Optional: Visualize the sample data
    plt.hist(sample_data, bins=10, alpha=0.7, color='blue',
             edgecolor='black')
    plt.axvline(population_mean, color='red', linestyle='dashed',
                linewidth=2, label='Population Mean')
    plt.axvline(np.mean(sample_data), color='green',
                linestyle='dashed', linewidth=2, label='Sample Mean')
    plt.title('Histogram of Sample Data')
    plt.xlabel('Value')
    plt.ylabel('Frequency')
    plt.legend()
    plt.show()

```


Sample Mean: 99.60

Population Mean: 100.00

Z-Score: -0.14

P-Value: 0.8849

Fail to reject the null hypothesis: There is no significant difference between the sample mean and the population mean.



Q3.Implement a one-sample Z-test using Python to compare the sample mean with the population mean?

Ans:-

```
import numpy as np
import scipy.stats as stats

# Function to perform a one-sample Z-test
def one_sample_z_test(sample, population_mean, population_std,
alpha=0.05):
    # Calculate sample mean and sample size
    sample_mean = np.mean(sample)
    n = len(sample)

    # Calculate the Z-score
    z_score = (sample_mean - population_mean) / (population_std /
np.sqrt(n))

    # Calculate the p-value (two-tailed)
    p_value = 2 * (1 - stats.norm.cdf(abs(z_score)))

    # Print results
    print(f"Sample Mean: {sample_mean:.2f}")
```

```

print(f"Population Mean: {population_mean:.2f}")
print(f"Z-Score: {z_score:.2f}")
print(f"P-Value: {p_value:.4f}")

# Interpret the results
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant
difference between the sample mean and the population mean.")
else:
    print("Fail to reject the null hypothesis: There is no
significant difference between the sample mean and the population
mean.")

# Example usage
if __name__ == "__main__":
    # Sample data (you can modify this data)
    sample_data = [102, 98, 101, 105, 99, 100, 97, 104, 103, 101]

    # Known population parameters
    population_mean = 100 # Known population mean
    population_std = 10   # Known population standard deviation

    # Perform the one-sample Z-test
    one_sample_z_test(sample_data, population_mean, population_std)

Sample Mean: 101.00
Population Mean: 100.00
Z-Score: 0.32
P-Value: 0.7518
Fail to reject the null hypothesis: There is no significant difference
between the sample mean and the population mean.

```

Q4.Perform a two-tailed Z-test using Python and visualize the decision region on a plot?

Ans:-

```

import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import norm

# Function to perform a two-tailed Z-test
def two_tailed_z_test(sample, population_mean, population_std,
alpha=0.05):
    # Calculate sample mean and sample size
    sample_mean = np.mean(sample)
    n = len(sample)

    # Calculate the Z-score
    z_score = (sample_mean - population_mean) / (population_std /
np.sqrt(n))

```

```

# Calculate critical values
critical_val1 = norm.ppf(1 - alpha / 2)
critical_val2 = norm.ppf(alpha / 2)

# Print results
print(f"Sample Mean: {sample_mean:.2f}")
print(f"Population Mean: {population_mean:.2f}")
print(f"Z-Score: {z_score:.2f}")

# Decision based on Z-score
if abs(z_score) > critical_val1:
    print("Reject the null hypothesis: There is a significant difference.")
else:
    print("Fail to reject the null hypothesis: There is no significant difference.")

# Visualization
x = np.linspace(-4, 4, 1000)
y = norm.pdf(x, 0, 1)

plt.plot(x, y, label='Standard Normal Distribution', color='blue')
plt.axvline(x=critical_val1, color='red', linestyle='--',
label='Critical Value (Right)')
plt.axvline(x=critical_val2, color='red', linestyle='--',
label='Critical Value (Left)')
plt.fill_between(x, y, where=(x <= critical_val2) | (x >=
critical_val1), color='red', alpha=0.5, label='Rejection Region')
plt.axvline(x=z_score, color='green', linestyle='-', label='Z-
Score')

plt.title('Two-Tailed Z-Test')
plt.xlabel('Z-Score')
plt.ylabel('Probability Density')
plt.legend()
plt.grid()
plt.show()


# Example usage
if __name__ == "__main__":
    # Sample data (you can modify this data)
    sample_data = [102, 98, 101, 105, 99, 100, 97, 104, 103, 101]

    # Known population parameters
    population_mean = 100 # Known population mean
    population_std = 10 # Known population standard deviation

    # Perform the two-tailed Z-test
    two_tailed_z_test(sample_data, population_mean, population_std)

```

Sample Mean: 101.00
Population Mean: 100.00
Z-Score: 0.32
Fail to reject the null hypothesis: There is no significant difference.



Q5.Create a Python function that calculates and visualizes Type 1 and Type 2 errors during hypothesis testing?

Ans:-

```
import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import norm

def visualize_errors(population_mean, population_std, sample_size,
alpha=0.05):
    # Define the null and alternative hypotheses
    null_hypothesis_mean = population_mean
    alternative_hypothesis_mean = population_mean + 1 # Shift the
mean for the alternative hypothesis

    # Calculate critical value for Type I error
```

```

critical_value = norm.ppf(1 - alpha)

# Generate x values for the normal distributions
x = np.linspace(population_mean - 4 * population_std,
population_mean + 4 * population_std, 1000)

# Calculate the probability density functions for the null and
alternative hypotheses
null_pdf = norm.pdf(x, null_hypothesis_mean, population_std /
np.sqrt(sample_size))
alternative_pdf = norm.pdf(x, alternative_hypothesis_mean,
population_std / np.sqrt(sample_size))

# Calculate Type I error area (alpha)
type_1_error_area = norm.cdf(critical_value)

# Calculate Type II error area (beta)
beta = norm.cdf(critical_value, loc=alternative_hypothesis_mean,
scale=population_std / np.sqrt(sample_size))

# Plotting
plt.figure(figsize=(12, 6))

# Plot null hypothesis distribution
plt.plot(x, null_pdf, label='Null Hypothesis Distribution',
color='blue')
plt.fill_between(x, null_pdf, where=(x <= critical_value),
color='red', alpha=0.5, label='Type I Error Area ( $\alpha$ )')

# Plot alternative hypothesis distribution
plt.plot(x, alternative_pdf, label='Alternative Hypothesis
Distribution', color='orange')
plt.fill_between(x, alternative_pdf, where=(x <= critical_value),
color='green', alpha=0.5, label='Type II Error Area ( $\beta$ )')

# Add critical value line
plt.axvline(x=critical_value, color='black', linestyle='--',
label='Critical Value')

# Add labels and legend
plt.title('Type I and Type II Errors in Hypothesis Testing')
plt.xlabel('Value')
plt.ylabel('Probability Density')
plt.legend()
plt.grid()
plt.show()

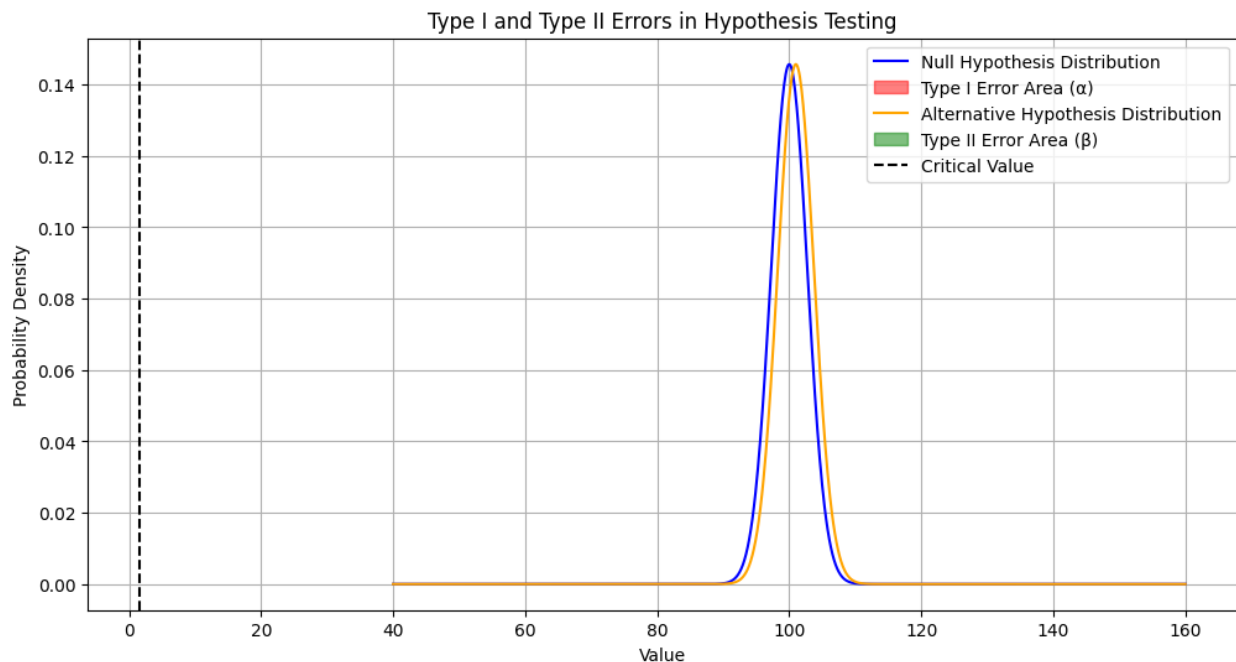
# Print Type I and Type II error probabilities
print(f"Type I Error ( $\alpha$ ): {type_1_error_area:.4f}")
print(f"Type II Error ( $\beta$ ): {beta:.4f}")

```



```
# Example usage
if __name__ == "__main__":
    population_mean = 100 # Known population mean
    population_std = 15 # Known population standard deviation
    sample_size = 30 # Sample size
    alpha = 0.05 # Significance level

    visualize_errors(population_mean, population_std, sample_size,
alpha)
```



Type I Error (α): 0.9500
Type II Error (β): 0.0000

Q6. Write a Python program to perform an independent T-test and interpret the results?

Ans:-

```
import numpy as np
import scipy.stats as stats

# Function to perform an independent two-sample T-test
def independent_t_test(group1, group2, alpha=0.05):
    # Calculate means and standard deviations
    mean1 = np.mean(group1)
    mean2 = np.mean(group2)
    std1 = np.std(group1, ddof=1) # Sample standard deviation
    std2 = np.std(group2, ddof=1) # Sample standard deviation
```

```

n1 = len(group1)
n2 = len(group2)

# Calculate the T-statistic and p-value
t_statistic, p_value = stats.ttest_ind(group1, group2)

# Print results
print(f"Group 1 Mean: {mean1:.2f}, Standard Deviation: {std1:.2f},
Sample Size: {n1}")
print(f"Group 2 Mean: {mean2:.2f}, Standard Deviation: {std2:.2f},
Sample Size: {n2}")
print(f"T-Statistic: {t_statistic:.2f}")
print(f"P-Value: {p_value:.4f}")

# Interpret the results
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant
difference between the two group means.")
else:
    print("Fail to reject the null hypothesis: There is no
significant difference between the two group means.")

# Example usage
if __name__ == "__main__":
    # Sample data for two independent groups
    group1 = [23, 21, 18, 25, 30, 22, 19, 24, 20, 26]
    group2 = [30, 32, 29, 35, 31, 28, 33, 34, 30, 29]

    # Perform the independent T-test
    independent_t_test(group1, group2)

Group 1 Mean: 22.80, Standard Deviation: 3.61, Sample Size: 10
Group 2 Mean: 31.10, Standard Deviation: 2.33, Sample Size: 10
T-Statistic: -6.10
P-Value: 0.0000
Reject the null hypothesis: There is a significant difference between
the two group means.

```

Q7.Perform a paired sample T-test using Python and visualize the comparison results?

Ans:-

```

import numpy as np
import scipy.stats as stats
import matplotlib.pyplot as plt

# Function to perform a paired sample T-test
def paired_t_test(group1, group2, alpha=0.05):
    # Calculate the T-statistic and p-value
    t_statistic, p_value = stats.ttest_rel(group1, group2)

```

```

# Print results
print(f"T-Statistic: {t_statistic:.2f}")
print(f"P-Value: {p_value:.4f}")

# Interpret the results
if p_value < alpha:
    print("Reject the null hypothesis: There is a significant
difference between the two paired samples.")
else:
    print("Fail to reject the null hypothesis: There is no
significant difference between the two paired samples.")

    return t_statistic, p_value

# Function to visualize the comparison results
def visualize_comparison(group1, group2):
    # Calculate means and standard deviations
    means = [np.mean(group1), np.mean(group2)]
    stds = [np.std(group1, ddof=1), np.std(group2, ddof=1)]
    labels = ['Group 1', 'Group 2']

    # Create a bar plot with error bars
    x = np.arange(len(labels))
    plt.bar(x, means, yerr=stds, capsize=5, color=['blue', 'orange'],
alpha=0.7)
    plt.xticks(x, labels)
    plt.ylabel('Mean Values')
    plt.title('Comparison of Paired Samples')
    plt.axhline(0, color='black', linewidth=0.8, linestyle='--')
    plt.grid(axis='y')
    plt.show()

# Example usage
if __name__ == "__main__":
    # Sample data for two paired groups
    group1 = [23, 21, 18, 25, 30, 22, 19, 24, 20, 26]
    group2 = [30, 32, 29, 35, 31, 28, 33, 34, 30, 29]

    # Perform the paired sample T-test
    t_statistic, p_value = paired_t_test(group1, group2)


    # Visualize the comparison results
    visualize_comparison(group1, group2)

```

T-Statistic: -6.56

P-Value: 0.0001

Reject the null hypothesis: There is a significant difference between the two paired samples.



Q8.Simulate data and perform both Z-test and T-test, then compare the results using Python?

Ans:-

```
import numpy as np
import scipy.stats as stats
import matplotlib.pyplot as plt

# Function to perform a one-sample Z-test
def z_test(sample, population_mean, population_std, alpha=0.05):
    sample_mean = np.mean(sample)
    n = len(sample)
    z_score = (sample_mean - population_mean) / (population_std /
np.sqrt(n))
    p_value = 2 * (1 - stats.norm.cdf(abs(z_score)))
    return z_score, p_value

# Function to perform a one-sample T-test
def t_test(sample, population_mean, alpha=0.05):
    t_statistic, p_value = stats.ttest_1samp(sample, population_mean)
    return t_statistic, p_value

# Function to simulate data
def simulate_data(sample_size, population_mean, population_std):
    return np.random.normal(loc=population_mean, scale=population_std,
```

```

size=sample_size)

# Main function to run the tests and compare results
def main():
    # Parameters for simulation
    sample_size = 30
    population_mean = 100
    population_std = 15

    # Simulate data
    sample_data = simulate_data(sample_size, population_mean,
                                population_std)

    # Perform Z-test
    z_score, z_p_value = z_test(sample_data, population_mean,
                                population_std)

    # Perform T-test
    t_statistic, t_p_value = t_test(sample_data, population_mean)

    # Print results
    print(f"Z-Test: Z-Score = {z_score:.2f}, P-Value = {z_p_value:.4f}")
    print(f"T-Test: T-Statistic = {t_statistic:.2f}, P-Value = {t_p_value:.4f}")

    # Visualization
    plt.figure(figsize=(10, 6))
    plt.hist(sample_data, bins=10, alpha=0.7, color='blue',
             edgecolor='black')
    plt.axvline(population_mean, color='red', linestyle='dashed',
                linewidth=2, label='Population Mean')
    plt.axvline(np.mean(sample_data), color='green',
                linestyle='dashed', linewidth=2, label='Sample Mean')
    plt.title('Histogram of Simulated Sample Data')
    plt.xlabel('Value')
    plt.ylabel('Frequency')
    plt.legend()
    plt.grid()
    plt.show()


# Run the main function
if __name__ == "__main__":
    main()

```

```

Z-Test: Z-Score = 0.42, P-Value = 0.6757
T-Test: T-Statistic = 0.54, P-Value = 0.5965

```



Q9. Write a Python function to calculate the confidence interval for a sample mean and explain its significance.

Ans:-

```
import numpy as np
import scipy.stats as stats

def confidence_interval(sample, confidence_level=0.95):
    """
    Calculate the confidence interval for the sample mean.

    Parameters:
    sample (array-like): The sample data.
    confidence_level (float): The confidence level (default is 0.95
    for 95% confidence).

    Returns:
    tuple: The lower and upper bounds of the confidence interval.
    """
    # Calculate sample mean and standard error
    sample_mean = np.mean(sample)
    sample_std = np.std(sample, ddof=1) # Sample standard deviation
    n = len(sample) # Sample size
    standard_error = sample_std / np.sqrt(n)
```

```

# Calculate the critical value for the t-distribution
critical_value = stats.t.ppf((1 + confidence_level) / 2, df=n-1)

# Calculate the margin of error
margin_of_error = critical_value * standard_error

# Calculate the confidence interval
lower_bound = sample_mean - margin_of_error
upper_bound = sample_mean + margin_of_error

return lower_bound, upper_bound

# Example usage
if __name__ == "__main__":
    # Sample data
    sample_data = [23, 21, 18, 25, 30, 22, 19, 24, 20, 26]

    # Calculate the confidence interval
    ci = confidence_interval(sample_data, confidence_level=0.95)
    print(f"95% Confidence Interval for the Sample Mean: {ci[0]:.2f}
to {ci[1]:.2f}")
95% Confidence Interval for the Sample Mean: 20.21 to 25.39

```

Q10. Write a Python program to calculate the margin of error for a given confidence level using sample data.

Ans:-

```

import numpy as np
import scipy.stats as stats

def margin_of_error(sample, confidence_level=0.95):
    """
    Calculate the margin of error for a given confidence level.

    Parameters:
    sample (array-like): The sample data.
    confidence_level (float): The confidence level (default is 0.95
for 95% confidence).

    Returns:
    float: The margin of error.
    """
    # Calculate sample standard deviation and sample size
    sample_std = np.std(sample, ddof=1) # Sample standard deviation
    n = len(sample) # Sample size

    # Calculate the standard error

```

```

standard_error = sample_std / np.sqrt(n)

# Calculate the critical value for the t-distribution
critical_value = stats.t.ppf((1 + confidence_level) / 2, df=n-1)

# Calculate the margin of error
margin_of_error = critical_value * standard_error

return margin_of_error

# Example usage
if __name__ == "__main__":
    # Sample data
    sample_data = [23, 21, 18, 25, 30, 22, 19, 24, 20, 26]

    # Calculate the margin of error
    me = margin_of_error(sample_data, confidence_level=0.95)
    print(f"Margin of Error at 95% Confidence Level: {me:.2f}")

```

Margin of Error at 95% Confidence Level: 2.59

Q11. Implement a Bayesian inference method using Bayes' Theorem in Python and explain the process.

Ans:-

```

import numpy as np
import matplotlib.pyplot as plt
from scipy.stats import beta

# Observed data
num_visitors = 1000 # Total number of visitors to the website
num_conversions = 50 # Number of conversions (desired actions)

# Prior hyperparameters for the Beta distribution
prior_alpha = 1 # Shape parameter
prior_beta = 1 # Shape parameter

# Update the prior with the observed data to get the posterior parameters
posterior_alpha = prior_alpha + num_conversions
posterior_beta = prior_beta + (num_visitors - num_conversions)

# Generate samples from the posterior Beta distribution
posterior_samples = np.random.beta(posterior_alpha, posterior_beta,
size=10000)

# Plot the posterior distribution
plt.figure(figsize=(8, 6))
plt.hist(posterior_samples, bins=30, density=True, color='skyblue',

```



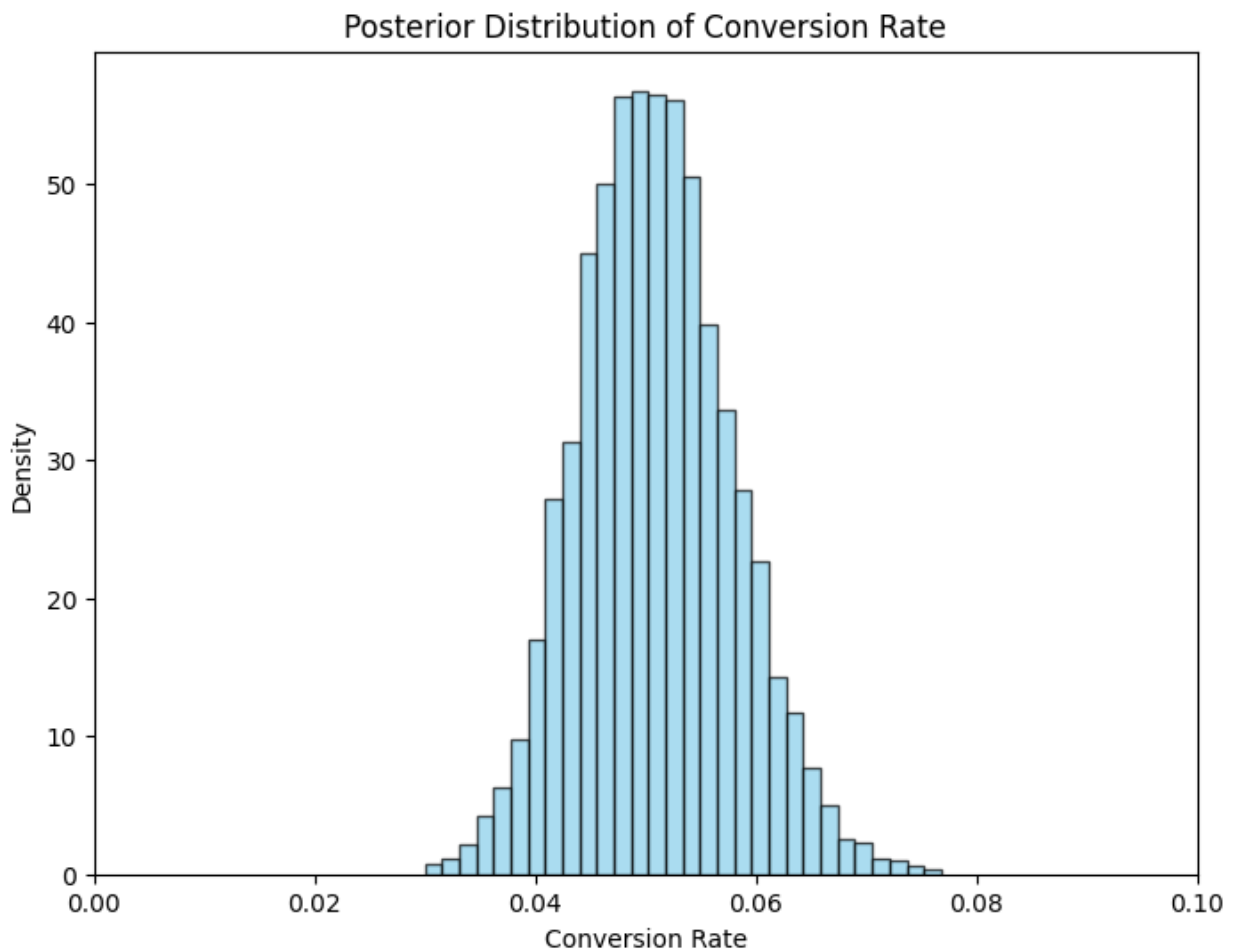
```

edgecolor='black', alpha=0.7)
plt.title('Posterior Distribution of Conversion Rate')
plt.xlabel('Conversion Rate')
plt.ylabel('Density')
plt.xlim(0, 0.1) # Limiting x-axis to focus on conversion rates close
to zero
plt.show()

# Calculate summary statistics
mean_conversion_rate = posterior_alpha / (posterior_alpha +
posterior_beta)
mode_conversion_rate = (posterior_alpha - 1) / (posterior_alpha +
posterior_beta - 2) # Mode of the Beta distribution

print("Mean conversion rate:", mean_conversion_rate)
print("Mode conversion rate:", mode_conversion_rate)

```



```

Mean conversion rate: 0.05089820359281437
Mode conversion rate: 0.05

```

Q12.Perform a Chi-square test for independence between two categorical variables in Python?

Ans:-

```
import numpy as np
import pandas as pd
import scipy.stats as stats

# Sample data: A contingency table
data = {
    'Category A': [30, 10, 20],
    'Category B': [20, 25, 15]
}

# Create a DataFrame
df = pd.DataFrame(data, index=['Group 1', 'Group 2', 'Group 3'])

# Display the contingency table
print("Contingency Table:")
print(df)

# Perform the Chi-square test
chi2_statistic, p_value, dof, expected = stats.chi2_contingency(df)

# Print the results
print("\nChi-square Statistic:", chi2_statistic)
print("P-value:", p_value)
print("Degrees of Freedom:", dof)
print("Expected Frequencies:")
print(expected)

# Interpret the results
alpha = 0.05
if p_value < alpha:
    print("\nReject the null hypothesis: There is a significant association between the two categorical variables.")
else:
    print("\nFail to reject the null hypothesis: There is no significant association between the two categorical variables.")
```

Contingency Table:

	Category A	Category B
Group 1	30	20
Group 2	10	25
Group 3	20	15

Chi-square Statistic: 9.142857142857144

P-value: 0.010343173196618245

Degrees of Freedom: 2

Expected Frequencies:

[[25. 25.]

```
[17.5 17.5]
[17.5 17.5]]
```

Reject the null hypothesis: There is a significant association between the two categorical variables.

Q13. Write a Python program to calculate the expected frequencies for a Chi-square test based on observed data?

Ans:-

```
import numpy as np
import pandas as pd

def calculate_expected_frequencies(observed):
    """
    Calculate the expected frequencies for a Chi-square test based on
    observed data.

    Parameters:
    observed (array-like): The observed data in a contingency table.

    Returns:
    np.ndarray: The expected frequencies.
    """
    # Convert observed data to a DataFrame if it's not already
    if not isinstance(observed, pd.DataFrame):
        observed = pd.DataFrame(observed)

    # Calculate row totals, column totals, and grand total
    row_totals = observed.sum(axis=1).values.reshape(-1, 1) # Reshape
    for broadcasting
    column_totals = observed.sum(axis=0).values # Column totals
    grand_total = observed.values.sum() # Grand total

    # Calculate expected frequencies
    expected = (row_totals * column_totals) / grand_total

    return expected

# Example usage
if __name__ == "__main__":
    # Sample observed data (contingency table)
    observed_data = np.array([[30, 10, 20],
                              [20, 25, 15]])

    # Create a DataFrame for better visualization
    observed_df = pd.DataFrame(observed_data, columns=['Category A',
    'Category B', 'Category C'], index=['Group 1', 'Group 2'])
```

```

print("Observed Frequencies:")
print(observed_df)

# Calculate expected frequencies
expected_frequencies =
calculate_expected_frequencies(observed_data)

# Display expected frequencies
expected_df = pd.DataFrame(expected_frequencies,
columns=['Category A', 'Category B', 'Category C'], index=['Group 1',
'Group 2'])
print("\nExpected Frequencies:")
print(expected_df)

```

Observed Frequencies:

	Category A	Category B	Category C
Group 1	30	10	20
Group 2	20	25	15

Expected Frequencies:

	Category A	Category B	Category C
Group 1	25.0	17.5	17.5
Group 2	25.0	17.5	17.5

Q14.Perform a goodness-of-fit test using Python to compare the observed data to an expected distribution.

Ans:-

```

import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt

def chi_square_goodness_of_fit(observed, expected):
    """
    Perform a Chi-square goodness-of-fit test.

    Parameters:
    observed (array-like): The observed frequencies.
    expected (array-like): The expected frequencies.

    Returns:
    chi2_statistic (float): The Chi-square statistic.
    p_value (float): The p-value of the test.
    """
    # Calculate the Chi-square statistic and p-value
    chi2_statistic, p_value = stats.chisquare(observed, expected)
    return chi2_statistic, p_value

```

```

# Example usage
if __name__ == "__main__":
    # Sample observed data (e.g., counts of outcomes)
    observed_data = np.array([50, 30, 20]) # Observed frequencies for
    three categories

    # Expected frequencies for a uniform distribution
    total_observed = observed_data.sum()
    num_categories = len(observed_data)
    expected_data = np.full(num_categories, total_observed /
    num_categories) # Equal expected frequencies

    # Perform the Chi-square goodness-of-fit test
    chi2_statistic, p_value =
    chi_square_goodness_of_fit(observed_data, expected_data)

    # Print the results
    print("Observed Frequencies:", observed_data)
    print("Expected Frequencies:", expected_data)
    print("Chi-square Statistic:", chi2_statistic)
    print("P-value:", p_value)

    # Interpret the results
    alpha = 0.05
    if p_value < alpha:
        print("\nReject the null hypothesis: The observed data does
    not fit the expected distribution.")
    else:
        print("\nFail to reject the null hypothesis: The observed data
    fits the expected distribution.")

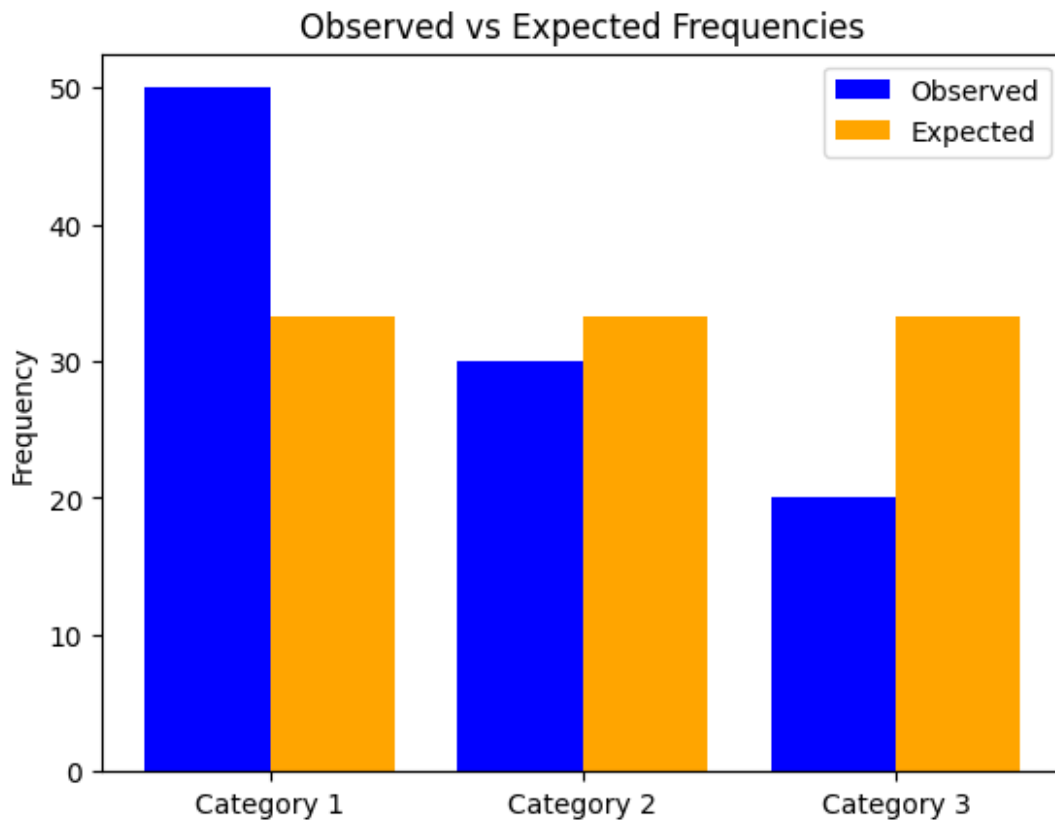
    # Optional: Visualize the observed vs expected frequencies
    categories = ['Category 1', 'Category 2', 'Category 3']
    x = np.arange(len(categories))

    plt.bar(x - 0.2, observed_data, width=0.4, label='Observed',
    color='blue')
    plt.bar(x + 0.2, expected_data, width=0.4, label='Expected',
    color='orange')
    plt.xticks(x, categories)
    plt.ylabel('Frequency')
    plt.title('Observed vs Expected Frequencies')
    plt.legend()
    plt.show()

Observed Frequencies: [50 30 20]
Expected Frequencies: [33.33333333 33.33333333 33.33333333]
Chi-square Statistic: 14.0
P-value: 0.0009118819655545164

```

Reject the null hypothesis: The observed data does not fit the expected distribution.



Q15. Create a Python script to simulate and visualize the Chi-square distribution and discuss its characteristics.

Ans:-

```
import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as stats

def plot_chi_square_distribution(degrees_of_freedom, x_range):
    """
    Plot the Chi-square distribution for given degrees of freedom.

    Parameters:
    degrees_of_freedom (list): List of degrees of freedom to plot.
    x_range (tuple): Range of x values for the plot.
    """
    x = np.linspace(x_range[0], x_range[1], 1000)
    plt.figure(figsize=(10, 6))
```

```

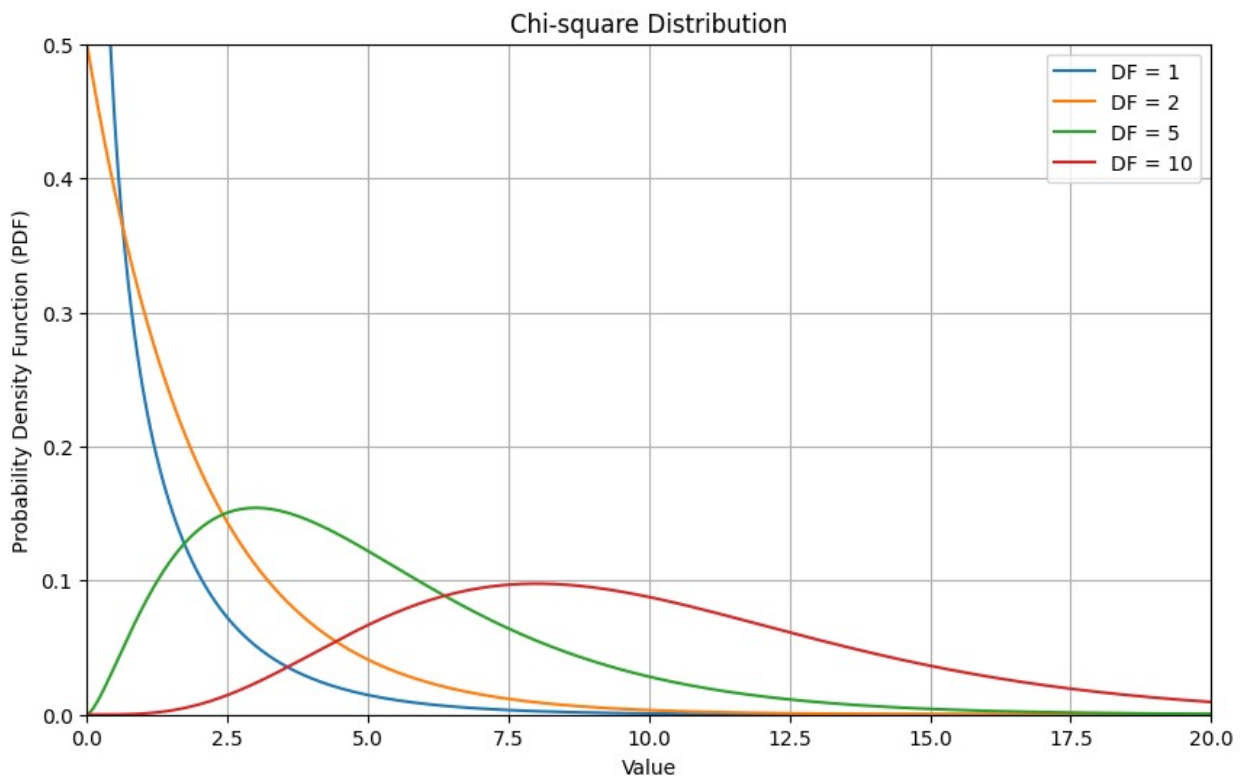
for df in degrees_of_freedom:
    # Calculate the Chi-square PDF
    pdf = stats.chi2.pdf(x, df)
    plt.plot(x, pdf, label=f'DF = {df}')

plt.title('Chi-square Distribution')
plt.xlabel('Value')
plt.ylabel('Probability Density Function (PDF)')
plt.legend()
plt.grid()
plt.xlim(x_range)
plt.ylim(0, 0.5)
plt.show()

# Example usage
if __name__ == "__main__":
    degrees_of_freedom = [1, 2, 5, 10] # Different degrees of freedom
    to visualize
    x_range = (0, 20) # Range of x values for the plot

    plot_chi_square_distribution(degrees_of_freedom, x_range)

```



Q16.Implement an F-test using Python to compare the variances of two random samples.

Ans:-

```

import numpy as np
import scipy.stats as stats

def f_test(sample1, sample2, alpha=0.05):
    """
    Perform an F-test to compare the variances of two samples.

    Parameters:
    sample1 (array-like): The first sample data.
    sample2 (array-like): The second sample data.
    alpha (float): The significance level (default is 0.05).

    Returns:
    f_statistic (float): The F-statistic.
    p_value (float): The p-value of the test.
    """
    # Calculate the variances of the samples
    var1 = np.var(sample1, ddof=1) # Sample variance
    var2 = np.var(sample2, ddof=1) # Sample variance

    # Calculate the F-statistic
    f_statistic = var1 / var2

    # Calculate the degrees of freedom
    df1 = len(sample1) - 1 # Degrees of freedom for sample1
    df2 = len(sample2) - 1 # Degrees of freedom for sample2

    # Calculate the p-value
    p_value = 1 - stats.f.cdf(f_statistic, df1, df2)

    return f_statistic, p_value

# Example usage
if __name__ == "__main__":
    # Sample data
    sample1 = np.array([20, 22, 19, 24, 21, 23, 20, 22, 19, 21])
    sample2 = np.array([30, 32, 29, 35, 31, 28, 33, 34, 30, 29])

    # Perform the F-test
    f_statistic, p_value = f_test(sample1, sample2)

    # Print the results
    print("F-Statistic:", f_statistic)
    print("P-Value:", p_value)

    # Interpret the results
    alpha = 0.05
    if p_value < alpha:
        print("\nReject the null hypothesis: The variances of the two samples are significantly different.")
    else:

```



```
print("\nFail to reject the null hypothesis: The variances of  
the two samples are not significantly different.")
```

F-Statistic: 0.5092024539877301

P-Value: 0.8354277292129488

Fail to reject the null hypothesis: The variances of the two samples are not significantly different.

Q17. Write a Python program to perform an ANOVA test to compare means between multiple groups and interpret the results.

Ans:-

```
import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt

def perform_anova(*groups):
    """
    Perform a one-way ANOVA test to compare means between multiple
    groups.

    Parameters:
    *groups: Variable number of group samples (arrays).

    Returns:
    f_statistic (float): The F-statistic.
    p_value (float): The p-value of the test.
    """
    # Perform the ANOVA test
    f_statistic, p_value = stats.f_oneway(*groups)
    return f_statistic, p_value

def plot_groups(groups):
    """
    Plot the boxplot of the groups for visual comparison.

    Parameters:
    groups: List of group samples (arrays).
    """
    plt.figure(figsize=(10, 6))
    plt.boxplot(groups, labels=[f'Group {i+1}' for i in
    range(len(groups))])
    plt.title('Boxplot of Groups')
    plt.ylabel('Values')
    plt.grid()
    plt.show()
```

```

# Example usage
if __name__ == "__main__":
    # Sample data for three groups
    group1 = np.array([23, 21, 18, 25, 30])
    group2 = np.array([30, 32, 29, 35, 31])
    group3 = np.array([22, 24, 20, 23, 21])

    # Perform the ANOVA test
    f_statistic, p_value = perform_anova(group1, group2, group3)

    # Print the results
    print("F-Statistic:", f_statistic)
    print("P-Value:", p_value)

    # Interpret the results
    alpha = 0.05
    if p_value < alpha:
        print("\nReject the null hypothesis: At least one group mean
is significantly different.")
    else:
        print("\nFail to reject the null hypothesis: There is no
significant difference between group means.")

    # Plot the groups for visual comparison
    plot_groups([group1, group2, group3])

```

F-Statistic: 13.729537366548039

P-Value: 0.0007910539293303842


Reject the null hypothesis: At least one group mean is significantly different.

<ipython-input-17-343a11efb66b>:29: MatplotlibDeprecationWarning: The 'labels' parameter of boxplot() has been renamed 'tick_labels' since Matplotlib 3.9; support for the old name will be dropped in 3.11.

```

plt.boxplot(groups, labels=[f'Group {i+1}' for i in
range(len(groups))])

```



Q18. Perform a one-way ANOVA test using Python to compare the means of different groups and plot the results.

Ans:-

```
import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt

def perform_anova(*groups):
    """
    Perform a one-way ANOVA test to compare means between multiple
    groups.

    Parameters:
    *groups: Variable number of group samples (arrays).

    Returns:
    f_statistic (float): The F-statistic.
    p_value (float): The p-value of the test.
    """
    # Perform the ANOVA test
    f_statistic, p_value = stats.f_oneway(*groups)
    return f_statistic, p_value

def plot_groups(groups, group_labels):
```

```

"""
Plot the boxplot of the groups for visual comparison.

Parameters:
groups: List of group samples (arrays).
group_labels: List of labels for each group.
"""

plt.figure(figsize=(10, 6))
plt.boxplot(groups, labels=group_labels)
plt.title('Boxplot of Groups')
plt.ylabel('Values')
plt.grid()
plt.show()

# Example usage
if __name__ == "__main__":
    # Sample data for three groups
    group1 = np.array([23, 21, 18, 25, 30])
    group2 = np.array([30, 32, 29, 35, 31])
    group3 = np.array([22, 24, 20, 23, 21])

    # Perform the ANOVA test
    f_statistic, p_value = perform_anova(group1, group2, group3)

    # Print the results
    print("F-Statistic:", f_statistic)
    print("P-Value:", p_value)

    # Interpret the results
    alpha = 0.05
    if p_value < alpha:
        print("\nReject the null hypothesis: At least one group mean
is significantly different.")
    else:
        print("\nFail to reject the null hypothesis: There is no
significant difference between group means.")

    # Plot the groups for visual comparison
    plot_groups([group1, group2, group3], ['Group 1', 'Group 2',
'Group 3'])


F-Statistic: 13.729537366548039
P-Value: 0.0007910539293303842

Reject the null hypothesis: At least one group mean is significantly
different.

<ipython-input-18-584463ee2571>:30: MatplotlibDeprecationWarning: The
'labels' parameter of boxplot() has been renamed 'tick_labels' since

```

Matplotlib 3.9; support for the old name will be dropped in 3.11.
`plt.boxplot(groups, labels=group_labels)`



Q19. Write a Python function to check the assumptions (normality, independence, and equal variance) for ANOVA.

Ans:-

```
import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns

def check_anova_assumptions(*groups):
    """
    Check the assumptions for ANOVA: normality, independence, and
    equal variance.

    Parameters:
    *groups: Variable number of group samples (arrays).
    """
    # Check Normality using Shapiro-Wilk test
    print("Normality Test Results:")
    for i, group in enumerate(groups):
        stat, p_value = stats.shapiro(group)
```

```

        print(f"Group {i+1}: W-statistic = {stat:.4f}, p-value =
{p_value:.4f}")
        if p_value < 0.05:
            print(f"Group {i+1} does not follow a normal distribution
(reject H0).\n")
        else:
            print(f"Group {i+1} follows a normal distribution (fail to
reject H0).\n")

    # Check Equal Variance using Levene's test
    stat, p_value = stats.levene(*groups)
    print("Equal Variance Test (Levene's Test):")
    print(f"Levene's W-statistic = {stat:.4f}, p-value =
{p_value:.4f}")
    if p_value < 0.05:
        print("The variances are significantly different (reject H0).\n")
    else:
        print("The variances are equal (fail to reject H0).\n")

    # Visualize the data distribution
    plt.figure(figsize=(12, 6))
    for i, group in enumerate(groups):
        sns.kdeplot(group, label=f'Group {i+1}', fill=True)
    plt.title('Kernel Density Estimation of Groups')
    plt.xlabel('Values')
    plt.ylabel('Density')
    plt.legend()
    plt.grid()
    plt.show()

    # Boxplot for visual inspection of equal variance
    plt.figure(figsize=(10, 6))
    plt.boxplot(groups, labels=[f'Group {i+1}' for i in
range(len(groups))])
    plt.title('Boxplot of Groups')
    plt.ylabel('Values')
    plt.grid()
    plt.show()

# Example usage
if __name__ == "__main__":
    # Sample data for three groups
    group1 = np.random.normal(loc=20, scale=5, size=30)
    group2 = np.random.normal(loc=22, scale=5, size=30)
    group3 = np.random.normal(loc=21, scale=5, size=30)

    # Check ANOVA assumptions
    check_anova_assumptions(group1, group2, group3)

```

Normality Test Results:

Group 1: W-statistic = 0.9838, p-value = 0.9157

Group 1 follows a normal distribution (fail to reject H_0).

Group 2: W-statistic = 0.9456, p-value = 0.1287

Group 2 follows a normal distribution (fail to reject H_0).


Group 3: W-statistic = 0.9666, p-value = 0.4510

Group 3 follows a normal distribution (fail to reject H_0).

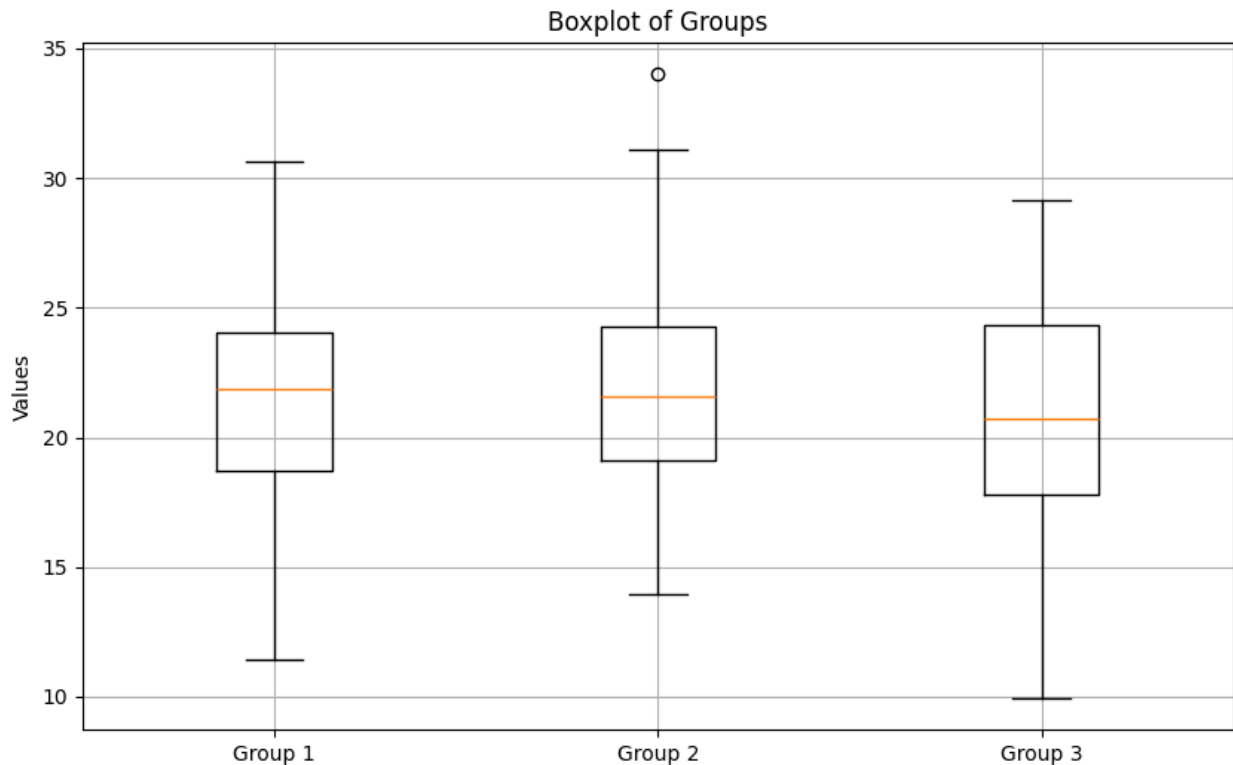
Equal Variance Test (Levene's Test):

Levene's W-statistic = 0.1427, p-value = 0.8672

The variances are equal (fail to reject H_0).



```
<ipython-input-19-4a75d65e6e31>:46: MatplotlibDeprecationWarning: The  
'labels' parameter of boxplot() has been renamed 'tick_labels' since  
Matplotlib 3.9; support for the old name will be dropped in 3.11.  
plt.boxplot(groups, labels=[f'Group {i+1}' for i in  
range(len(groups))])
```



Q20. Perform a two-way ANOVA test using Python to study the interaction between two factors and visualize the results.

Ans:- Step 1: Import Required Libraries

python

```
import pandas as pd
import numpy as np
import statsmodels.api as sm
import statsmodels.formula.api as smf
import seaborn as sns
import matplotlib.pyplot as plt
```

Step 2: Create or Load Your Data

For demonstration, let's create a sample dataset. In practice, you would load your data from a CSV or other source.

```
# Create a sample dataset
np.random.seed(42)
data = {
    'Factor1': np.repeat(['A', 'B', 'C'], 20),
    'Factor2': np.tile(np.repeat(['X', 'Y'], 10), 3),
    'Response': np.random.normal(loc=0, scale=1, size=60)
```



```
}
df = pd.DataFrame(data)
```

Step 3: Perform Two-Way ANOVA

```
# Fit the model
model = smf.ols('Response ~ C(Factor1) * C(Factor2)', data=df).fit()

# Perform ANOVA
anova_table = sm.stats.anova_lm(model, typ=2)
print(anova_table)
```

	sum_sq	df	F	PR(>F)
C(Factor1)	0.580882	2.0	0.398223	0.673468
C(Factor2)	1.275222	1.0	1.748451	0.191645
C(Factor1):C(Factor2)	7.458570	2.0	5.113208	0.009255
Residual	39.384553	54.0	NaN	NaN

Step 4: Check for Significant Interactions

The ANOVA table will show you the p-values for the main effects and their interaction. If the p-value for the interaction term is less than 0.05, it indicates a significant interaction between the two factors.

Step 5: Visualize the Results

You can visualize the interaction using an interaction plot.

```
import pandas as pd
import numpy as np
import statsmodels.api as sm
import statsmodels.formula.api as smf
import seaborn as sns
import matplotlib.pyplot as plt


# Create a sample dataset
np.random.seed(42)
data = {
    'Factor1': np.repeat(['A', 'B', 'C'], 20),
    'Factor2': np.tile(np.repeat(['X', 'Y'], 10), 3),
    'Response': np.random.normal(loc=0, scale=1, size=60)
}
df = pd.DataFrame(data)

# Fit the model
model = smf.ols('Response ~ C(Factor1) * C(Factor2)', data=df).fit()

# Perform ANOVA
anova_table = sm.stats.anova_lm(model, typ=2)
print(anova_table)
```

```
# Create interaction plot using lineplot
plt.figure(figsize=(8, 6))
sns.lineplot(data=df, x='Factor1', y='Response', hue='Factor2',
             estimator='mean', marker='o')
plt.title('Interaction Plot')
plt.xlabel('Factor 1')
plt.ylabel('Mean Response')
plt.grid()
plt.show()
```


	sum_sq	df	F	PR(>F)
C(Factor1)	0.580882	2.0	0.398223	0.673468
C(Factor2)	1.275222	1.0	1.748451	0.191645
C(Factor1):C(Factor2)	7.458570	2.0	5.113208	0.009255
Residual	39.384553	54.0	NaN	NaN



Step 6: Boxplot for Group Comparison

Additionally, you can create boxplots to visualize the distribution of the response variable across the different groups.

```
# Boxplot
plt.figure(figsize=(10, 6))
sns.boxplot(x='Factor1', y='Response', hue='Factor2', data=df)
plt.title('Boxplot of Response by Factor1 and Factor2')
plt.xlabel('Factor 1')
plt.ylabel('Response')
plt.legend(title='Factor 2')
plt.show()
```



Q21. Write a Python program to visualize the F-distribution and discuss its use in hypothesis testing.

Ans:-

```
import numpy as np
import matplotlib.pyplot as plt
import scipy.stats as stats

def plot_f_distribution(degrees_of_freedom1, degrees_of_freedom2):
    """
    Plot the F-distribution for given degrees of freedom.

    Parameters:
    degrees_of_freedom1 (int): Degrees of freedom for the numerator.
    degrees_of_freedom2 (int): Degrees of freedom for the denominator.
```

```


"""
x = np.linspace(0, 5, 1000) # Range of x values
pdf = stats.f.pdf(x, degrees_of_freedom1, degrees_of_freedom2) #
F-distribution PDF

plt.figure(figsize=(10, 6))
plt.plot(x, pdf, label=f'F-distribution
(df1={degrees_of_freedom1}, df2={degrees_of_freedom2})', color='blue')
plt.title('F-Distribution')
plt.xlabel('F-value')
plt.ylabel('Probability Density Function (PDF)')
plt.grid()
plt.legend()
plt.xlim(0, 5)
plt.ylim(0, 1)
plt.show()

# Example usage
if __name__ == "__main__":
    # Define degrees of freedom
    df1 = 5 # Degrees of freedom for the numerator
    df2 = 2 # Degrees of freedom for the denominator

    # Plot the F-distribution
    plot_f_distribution(df1, df2)

```



Q22.Perform a one-way ANOVA test in Python and visualize the results with boxplots to compare group means.

Ans:-

```
import numpy as np
import pandas as pd
import scipy.stats as stats
import seaborn as sns
import matplotlib.pyplot as plt

# Create a sample dataset
np.random.seed(42)
group1 = np.random.normal(loc=20, scale=5, size=30) # Group 1
group2 = np.random.normal(loc=22, scale=5, size=30) # Group 2
group3 = np.random.normal(loc=25, scale=5, size=30) # Group 3

# Combine the groups into a DataFrame
data = {
    'Group': ['Group 1'] * 30 + ['Group 2'] * 30 + ['Group 3'] * 30,
    'Response': np.concatenate([group1, group2, group3])
}
df = pd.DataFrame(data)

# Perform the one-way ANOVA test
f_statistic, p_value = stats.f_oneway(df[df['Group'] == 'Group 1']
['Response'],
df[df['Group'] == 'Group 2']
['Response'],
df[df['Group'] == 'Group 3']
['Response'])

# Print the results
print("F-Statistic:", f_statistic)
print("P-Value:", p_value)

# Interpret the results
alpha = 0.05
if p_value < alpha:
    print("\nReject the null hypothesis: At least one group mean is
significantly different.")
else:
    print("\nFail to reject the null hypothesis: There is no
significant difference between group means.")

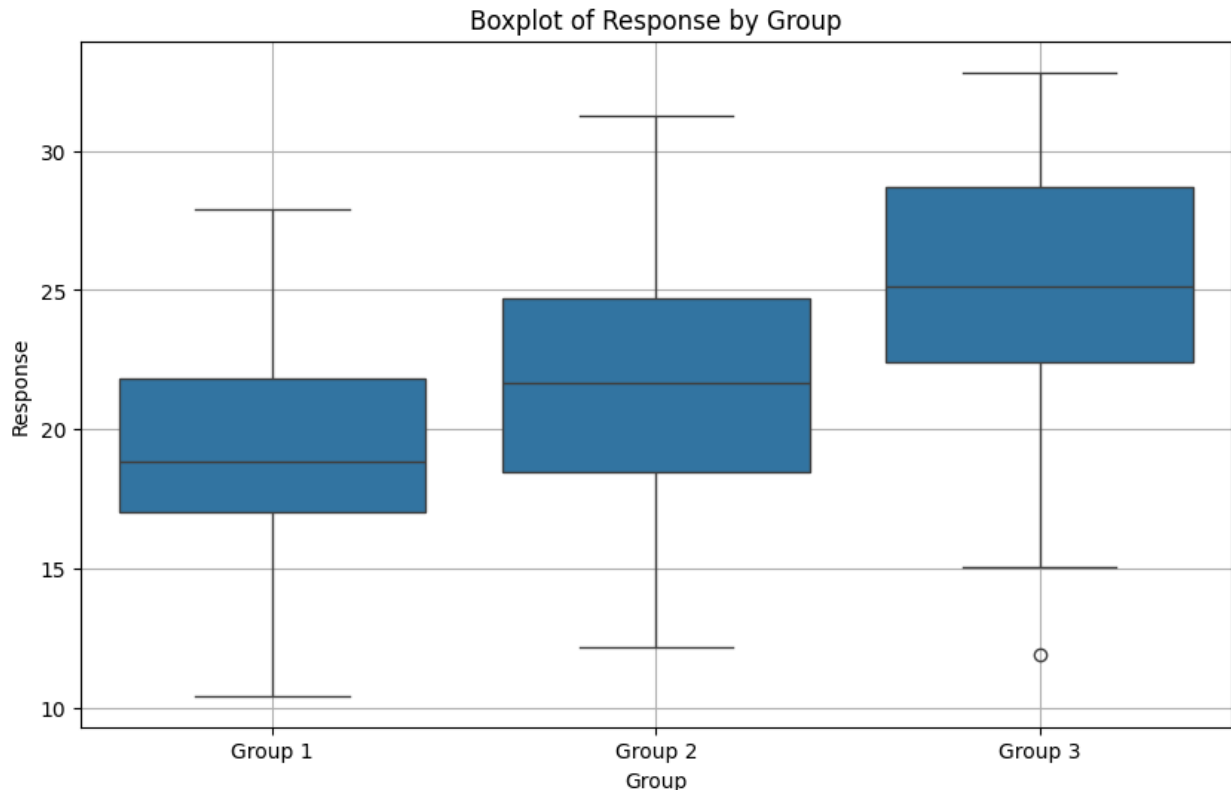
# Visualize the results with boxplots
plt.figure(figsize=(10, 6))
sns.boxplot(x='Group', y='Response', data=df)
plt.title('Boxplot of Response by Group')
plt.xlabel('Group')
```

```
plt.ylabel('Response')
plt.grid()
plt.show()
```

F-Statistic: 12.397871901825155

P-Value: 1.8305780660638664e-05

Reject the null hypothesis: At least one group mean is significantly different.



Q23. Simulate random data from a normal distribution, then perform hypothesis testing to evaluate the means.

Ans:- Python Code for Simulating Data and Performing Hypothesis Testing

```
import numpy as np
import scipy.stats as stats
import matplotlib.pyplot as plt

# Set parameters for the normal distribution
np.random.seed(42) # For reproducibility
mean = 50           # True mean of the population
std_dev = 10        # Standard deviation of the population
sample_size = 100   # Number of samples to simulate
```

```

# Simulate random data from a normal distribution
data = np.random.normal(loc=mean, scale=std_dev, size=sample_size)

# Perform a one-sample t-test
population_mean = 52 # Hypothesized population mean
t_statistic, p_value = stats.ttest_1samp(data, population_mean)

# Print the results
print("Sample Mean:", np.mean(data))
print("T-Statistic:", t_statistic)
print("P-Value:", p_value)

# Interpret the results
alpha = 0.05
if p_value < alpha:
    print("\nReject the null hypothesis: The sample mean is
significantly different from the population mean.")
else:
    print("\nFail to reject the null hypothesis: The sample mean is
not significantly different from the population mean.")

# Visualize the data
plt.figure(figsize=(10, 6))
plt.hist(data, bins=15, alpha=0.7, color='blue', edgecolor='black')
plt.axvline(population_mean, color='red', linestyle='dashed',
linewidth=2, label='Hypothesized Mean')
plt.axvline(np.mean(data), color='green', linestyle='dashed',
linewidth=2, label='Sample Mean')
plt.title('Histogram of Simulated Data')
plt.xlabel('Value')
plt.ylabel('Frequency')
plt.legend()
plt.grid()
plt.show()

```

```

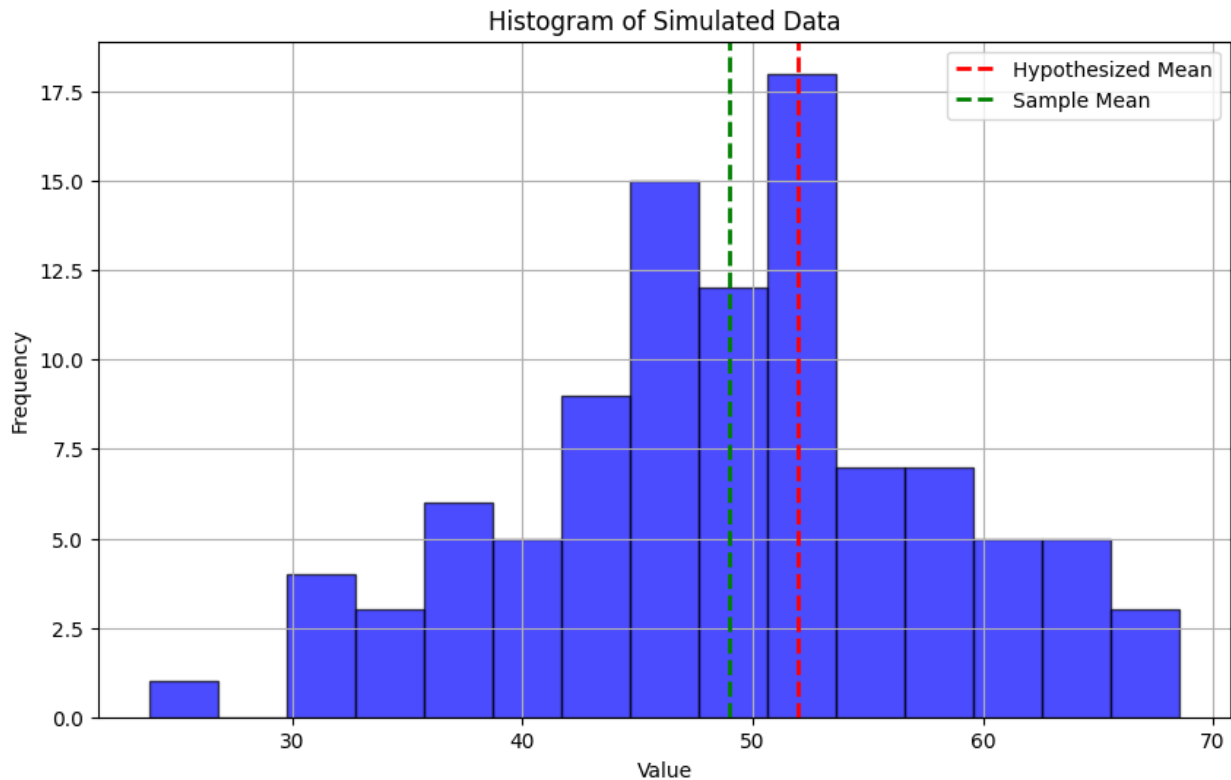
Sample Mean: 48.96153482605907
T-Statistic: -3.3457066775666733
P-Value: 0.0011606436334217142

```

```

Reject the null hypothesis: The sample mean is significantly different
from the population mean.

```



Q24. Perform a hypothesis test for population variance using a Chi-square distribution and interpret the results.

Ans:-

```
import numpy as np
import scipy.stats as stats

# Sample data
np.random.seed(42) # For reproducibility
data = np.random.normal(loc=50, scale=10, size=30) # Simulated data

# Hypothesized population variance
sigma_squared_0 = 100 # Hypothesized variance (10^2)

# Calculate sample variance and sample size
n = len(data)
sample_variance = np.var(data, ddof=1) # Sample variance with Bessel's correction

# Calculate the Chi-square statistic
chi_square_statistic = (n - 1) * sample_variance / sigma_squared_0

# Degrees of freedom
df = n - 1
```



```

# Significance level
alpha = 0.05

# Critical values for two-tailed test
critical_value_lower = stats.chi2.ppf(alpha / 2, df)
critical_value_upper = stats.chi2.ppf(1 - alpha / 2, df)

# Print results
print("Sample Variance:", sample_variance)
print("Chi-square Statistic:", chi_square_statistic)
print("Critical Value Lower:", critical_value_lower)
print("Critical Value Upper:", critical_value_upper)

# Interpret the results
if chi_square_statistic < critical_value_lower or chi_square_statistic
> critical_value_upper:
    print("\nReject the null hypothesis: The population variance is
significantly different from the hypothesized variance.")
else:
    print("\nFail to reject the null hypothesis: The population
variance is not significantly different from the hypothesized
variance.")

Sample Variance: 81.00115693656518
Chi-square Statistic: 23.490335511603902
Critical Value Lower: 16.04707169536489
Critical Value Upper: 45.72228580417452

Fail to reject the null hypothesis: The population variance is not
significantly different from the hypothesized variance.

```

Q25. Write a Python script to perform a Z-test for comparing proportions between two datasets or groups.

Ans:-

```

import numpy as np
import scipy.stats as stats

def z_test_proportions(success_a, n_a, success_b, n_b):
    """
    Perform a Z-test for comparing proportions between two groups.

    Parameters:
    success_a (int): Number of successes in group A.
    n_a (int): Total number of observations in group A.
    success_b (int): Number of successes in group B.
    n_b (int): Total number of observations in group B.

    Returns:

```

```

z_statistic (float): The Z-statistic.
p_value (float): The p-value of the test.
"""

# Calculate proportions
p_a = success_a / n_a
p_b = success_b / n_b

# Calculate the pooled proportion
p_pool = (success_a + success_b) / (n_a + n_b)

# Calculate the standard error
standard_error = np.sqrt(p_pool * (1 - p_pool) * (1/n_a + 1/n_b))

# Calculate the Z-statistic
z_statistic = (p_a - p_b) / standard_error

# Calculate the p-value (two-tailed)
p_value = 2 * (1 - stats.norm.cdf(abs(z_statistic)))

return z_statistic, p_value

# Example usage
if __name__ == "__main__":
    # Data for group A
    success_a = 30 # Number of successes in group A
    n_a = 100      # Total observations in group A

    # Data for group B
    success_b = 20 # Number of successes in group B
    n_b = 80      # Total observations in group B

    # Perform the Z-test for proportions
    z_statistic, p_value = z_test_proportions(success_a, n_a,
success_b, n_b)

    # Print the results
    print("Z-Statistic:", z_statistic)
    print("P-Value:", p_value)

    # Interpret the results
    alpha = 0.05
    if p_value < alpha:
        print("\nReject the null hypothesis: There is a significant
difference between the proportions.")
    else:
        print("\nFail to reject the null hypothesis: There is no
significant difference between the proportions.")

Z-Statistic: 0.7442084075352506
P-Value: 0.45675040242316944

```

Fail to reject the null hypothesis: There is no significant difference between the proportions.

Q26. Implement an F-test for comparing the variances of two datasets, then interpret and visualize the results.

Ans:-

```
import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt
import seaborn as sns

def f_test(sample1, sample2):
    """
    Perform an F-test to compare the variances of two samples.

    Parameters:
    sample1 (array-like): The first sample data.
    sample2 (array-like): The second sample data.

    Returns:
    f_statistic (float): The F-statistic.
    p_value (float): The p-value of the test.
    """
    # Calculate the variances of the samples
    var1 = np.var(sample1, ddof=1) # Sample variance
    var2 = np.var(sample2, ddof=1) # Sample variance

    # Calculate the F-statistic
    f_statistic = var1 / var2

    # Calculate the degrees of freedom
    df1 = len(sample1) - 1 # Degrees of freedom for sample1
    df2 = len(sample2) - 1 # Degrees of freedom for sample2

    # Calculate the p-value
    p_value = 1 - stats.f.cdf(f_statistic, df1, df2)

    return f_statistic, p_value

# Example usage
if __name__ == "__main__":
    # Generate sample data
    np.random.seed(42) # For reproducibility
    sample1 = np.random.normal(loc=50, scale=10, size=30) # Sample 1
    sample2 = np.random.normal(loc=55, scale=15, size=30) # Sample 2

    # Perform the F-test
```

```

f_statistic, p_value = f_test(sample1, sample2)

# Print the results
print("F-Statistic:", f_statistic)
print("P-Value:", p_value)

# Interpret the results
alpha = 0.05
if p_value < alpha:
    print("\nReject the null hypothesis: The variances of the two
samples are significantly different.")
else:
    print("\nFail to reject the null hypothesis: The variances of
the two samples are not significantly different.")

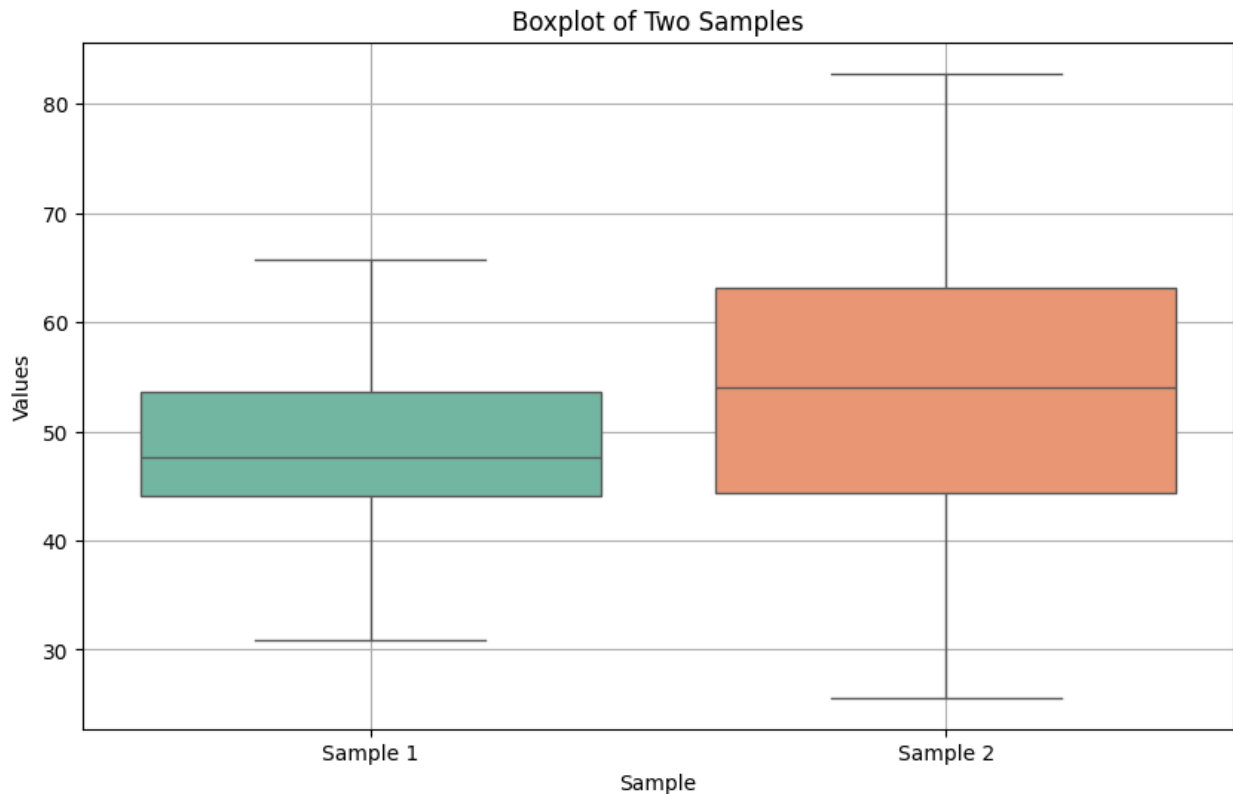
# Visualize the results with boxplots
plt.figure(figsize=(10, 6))
sns.boxplot(data=[sample1, sample2], palette="Set2")
plt.title('Boxplot of Two Samples')
plt.xlabel('Sample')
plt.ylabel('Values')
plt.xticks([0, 1], ['Sample 1', 'Sample 2'])
plt.grid()
plt.show()

```

F-Statistic: 0.41525416799961246

P-Value: 0.9895622022746069

Fail to reject the null hypothesis: The variances of the two samples are not significantly different.



Q27.Perform a Chi-square test for goodness of fit with simulated data and analyze the results.

Ans:-Python Code for Chi-Square Goodness of Fit Test

```
import numpy as np
import pandas as pd
import scipy.stats as stats
import matplotlib.pyplot as plt

def chi_square_goodness_of_fit(observed, expected):
    """
    Perform a Chi-square goodness of fit test.

    Parameters:
    observed (array-like): Observed frequencies.
    expected (array-like): Expected frequencies.

    Returns:
    chi2_statistic (float): The Chi-square statistic.
    p_value (float): The p-value of the test.
    """
    chi2_statistic, p_value = stats.chisquare(f_obs=observed,
f_exp=expected)
    return chi2_statistic, p_value

# Simulate categorical data
```

```

np.random.seed(42) # For reproducibility
categories = ['A', 'B', 'C', 'D']
observed_counts = np.random.randint(10, 30, size=len(categories)) #
Simulated observed counts

# Define expected proportions (e.g., equal distribution)
expected_proportions = [0.25, 0.25, 0.25, 0.25] # Equal proportions
for 4 categories
expected_counts = np.array(expected_proportions) *
sum(observed_counts) # Calculate expected counts

# Perform the Chi-square goodness of fit test
chi2_statistic, p_value = chi_square_goodness_of_fit(observed_counts,
expected_counts)

# Print the results
print("Observed Counts:", observed_counts)
print("Expected Counts:", expected_counts)
print("Chi-square Statistic:", chi2_statistic)
print("P-Value:", p_value)

# Interpret the results
alpha = 0.05
if p_value < alpha:
    print("\nReject the null hypothesis: The observed frequencies
differ significantly from the expected frequencies.")
else:
    print("\nFail to reject the null hypothesis: The observed
frequencies do not differ significantly from the expected
frequencies.")

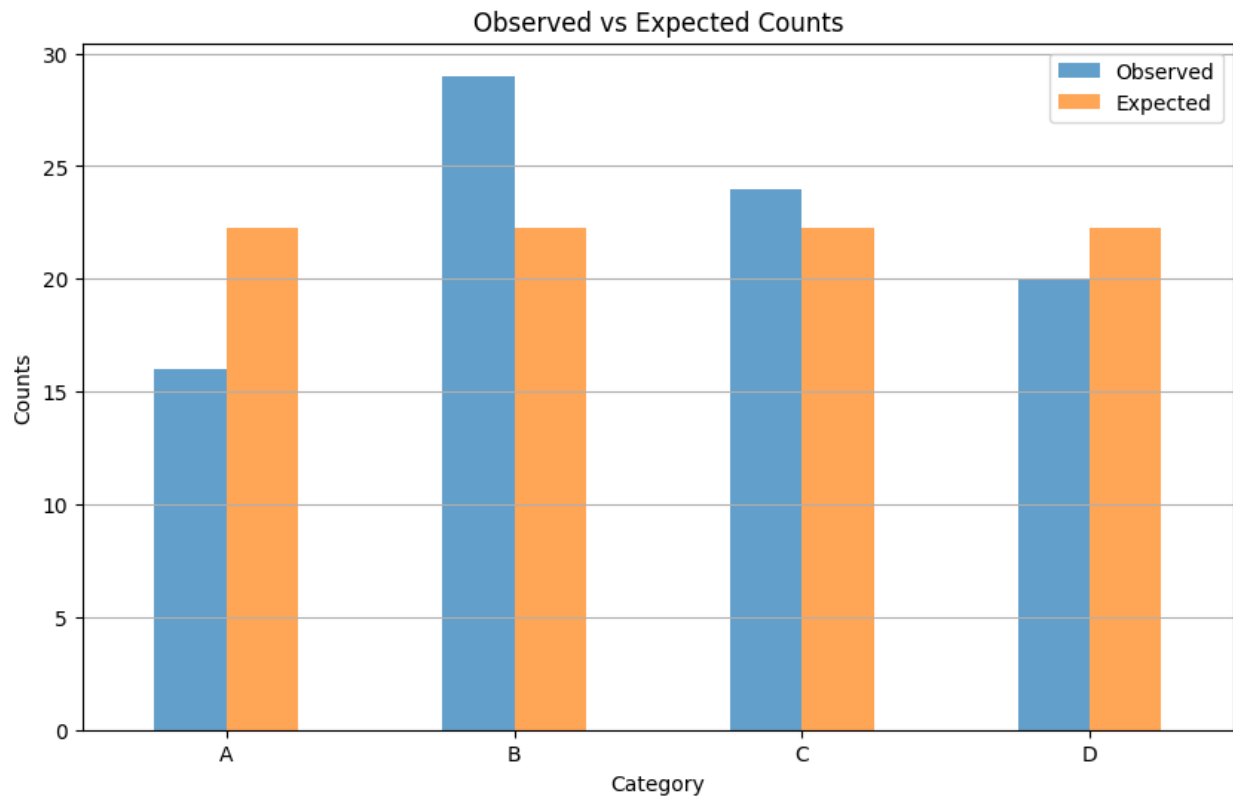
# Visualize the results
df = pd.DataFrame({'Category': categories, 'Observed':
observed_counts, 'Expected': expected_counts})
df.set_index('Category', inplace=True)

# Plotting
df.plot(kind='bar', figsize=(10, 6), alpha=0.7)
plt.title('Observed vs Expected Counts')
plt.ylabel('Counts')
plt.xticks(rotation=0)
plt.grid(axis='y')
plt.show()

Observed Counts: [16 29 24 20]
Expected Counts: [22.25 22.25 22.25 22.25]
Chi-square Statistic: 4.168539325842697
P-Value: 0.24383063094585364

```

Fail to reject the null hypothesis: The observed frequencies do not differ significantly from the expected frequencies.



Theory Question Answer Assignment :- <https://drive.google.com/file/d/1RCj49Bgsz5JHPqj-0qmjllasw18aj4r8/view?usp=drivesdk>